

Virtual Cell Models Inflate Perturbation Effect Sizes and Undermine Causal Gene Regulatory Network Recovery

Anonymous Authors¹

Abstract

Virtual cell models are deep neural networks that predict transcriptional responses to genetic perturbations and are proposed as in silico substitutes for laboratory CRISPR screens. Ahlmann-Eltze & Huber (2025) recently showed that current models fail to outperform simple linear baselines at per-gene prediction; we ask whether the same holds downstream when their predictions are used to recover causal gene regulatory networks. We introduce PerturbCausal, a benchmark that pushes predictions from four virtual cell models (GEARS, CPA, Geneformer, and the 2025 Arc Institute STATE foundation model) through four causal discovery algorithms (GIES, inspre, PC, NOTEARS) on the Replogle K562 and RPE1 Perturb-seq datasets and the Norman 2019 K562 dataset. Across five seeds, the three pre-2025 models do not significantly outperform a sparse linear regression trained on control cells alone. STATE appears to exceed the linear baseline at a single seed ($F_1 = 0.462$, 320-cell embedding subset), but this does not replicate: across three seeds $F_1 = 0.184$ with a 95% CI of $[-0.088, 0.457]$ that spans zero (§5), so even the 2025 foundation model does not robustly beat sparse linear regression, and its absolute causal F1 remains below 0.5. We reconcile two competing narratives in the recent literature, mode collapse (Mejia et al., 2025) and effect-size inflation, as two views of the same failure: deep models predict similar dense response patterns across perturbations. A model-agnostic quantile-matching calibration corrects this magnitude

failure exactly but recovers only about 40% of the causal gap, locating the remainder in joint-distribution structure that distributional calibration cannot reach. The benchmark, post-hoc quantile-matching calibration module, and reproducibility code are released as an open-source Python package.

1. Introduction

Virtual cell models (VCMs) predict cellular responses to genetic perturbations and are proposed as computational substitutes for CRISPR screens that cost millions of dollars per cell line (Replogle et al., 2022; Bunne et al., 2024). The dominant evaluation paradigm—exemplified by PerturbBench (Wu et al., 2024), the Virtual Cell Challenge (Roohani et al., 2025), and the Ahlmann-Eltze & Huber comparison (Ahlmann-Eltze & Huber, 2025)—measures per-gene reconstruction accuracy of predicted transcriptional responses. Across these benchmarks, deep learning methods including graph neural networks (GEARS; Roohani et al. 2024), conditional variational autoencoders (CPA; Lotfollahi et al. 2023), and pretrained foundation transformers (Geneformer, scGPT; Theodoris et al. 2023; Cui et al. 2024) consistently fail to outperform simple linear baselines on the marginal-prediction task.

Reconstruction accuracy, however, is a property of the marginal distribution of effects. The downstream applications that motivate VCM development—combination perturbation prediction, drug target identification, and gene regulatory network (GRN) inference—depend on the joint distribution of effects across genes (Chevalley et al., 2025; Callahan et al., 2025). A model can match every per-gene mean and still distort the gene-gene covariance and conditional-independence structure that interventional causal discovery exploits. CausalBench evaluates causal discovery algorithms on real Perturb-seq data (Chevalley et al., 2025) but does not test whether VCM predictions can substitute for that real data in the same pipeline.

¹Anonymous Institution, Anonymous City, Anonymous Region, Anonymous Country. AUTHORERR: Missing \icmlcorrespondingauthor.

Preliminary work. Under review by the International Conference on Machine Learning (ICML). Do not distribute.

A NeurIPS 2025 perspective paper called explicitly for causal-recovery evaluation of VCMs but did not implement one (Callahan et al., 2025).

Adjacent recent work has surfaced relevant pieces. Mejia et al. (2025) identified mode collapse in scRNA-seq perturbation models—predictions reverting to the dataset mean—and proposed weighted training losses. Miller et al. (2025) argued that conventional perturbation prediction metrics are poorly calibrated and proposed metric calibration via interpolated-duplicate baselines. Brown et al. (2025) introduced inspre, a sparse inverse regression algorithm for large-scale causal discovery from interventional data. None of these works has integrated VCM predictions into a downstream causal discovery pipeline to test the substitutability hypothesis directly.

We close this gap. The contributions of this paper are:

1. A causal-recovery benchmark for VCMs. Predictions from three architecturally distinct VCMs are passed through identical GIES (Hauser & Bühlmann, 2012) and inspre (Brown et al., 2025) pipelines on Replogle K562 and RPE1 Perturb-seq (Replogle et al., 2022), and resulting graphs are compared against real-data ground truth.
2. A diagnosis of the failure mode. All three deep models inflate predicted effect-size distributions by a factor of approximately 23 relative to real biology. The inflation magnitudes are identical to within 0.3 percentage points across radically different architectures (GNN, VAE, transformer), arguing the failure is objective-driven (per-gene MSE) rather than architecture-driven.
3. An empirical signal probe. The inspre algorithm exposes that CPA and Geneformer’s average causal effect (ACE) matrices fall below the regularization-path detection threshold, recovering zero edges where real data and ElasticNet recover approximately 12,000.
4. A diagnostic decomposition of the substitutability gap. Post-hoc quantile-matching calibration (Bolstad et al., 2003; Reinhard et al., 2001) corrects the magnitude marginal exactly and recovers a small share of the causal-F1 gap, demonstrating that the residual gap lives in joint-distribution structure (gene-gene covariance, edge orientation) that distributional calibration cannot reach. The decomposition is the contribution; the calibration itself is a standard technique.

2. Related Work

Causal discovery on Perturb-seq. GIES (Hauser & Bühlmann, 2012) is the canonical interventional causal discovery algorithm and the primary backend in CausalBench (Chevalley et al., 2025). inspre (Brown et al., 2025) operates on the average causal effect matrix and scales to thousands of nodes. Both are used here as parallel backends.

Perturbation prediction benchmarks. A growing 2024–2025 literature converges on a single conclusion: deep learning does not yet beat simple baselines on per-gene reconstruction (Ahlmann-Eltze & Huber, 2025; Wong et al., 2025; Csendes et al., 2025; Wei et al., 2025; Viñas Torné et al., 2025; Bendidi et al., 2024; Wenteler et al., 2025; Zhu et al., 2025; Li et al., 2025; Mejia et al., 2025). PerturbBench (Wu et al., 2024) and the Virtual Cell Challenge (Roohani et al., 2025) provide the field’s standard evaluation harnesses; Viñas Torné et al. (2025) introduce “systematic variation” as a confounder of Δ -based metrics and Wei et al. (2025) extend the comparison to 27 methods on 29 datasets. None of these evaluate downstream causal recovery from VCM predictions.

Causal discovery from Perturb-seq. A parallel literature builds causal-discovery methods directly on real Perturb-seq data: GIES (Hauser & Bühlmann, 2012), inspre (Brown et al., 2025), BEELINE (Pratapa et al., 2019) as a benchmark of GRN inference algorithms, Scribe (Qiu et al., 2020), CASCADE (Cao et al., 2025) for causality-oriented reprogramming design, and confounder-robust IV-based methods (Park et al., 2026). Kernfeld et al. (2024) argue transcriptome data alone cannot control false discoveries in causal network inference. PerturbCausal is downstream of all of these: it asks whether predicted perturbation data can substitute for the real data those methods consume.

Causal evaluation of VCMs. Callahan et al. (2025) proposed a causal-evaluation taxonomy without implementation. Wu et al. (2025) used causal discovery as a featurizer for perturbation target prediction—orthogonal direction to ours.

Calibration in perturbation prediction. Mejia et al. (2025) propose a weighted training loss that requires retraining. Miller et al. (2025) propose metric calibration that does not change model outputs and argue that conventional metrics underestimate deep-learning performance under their dynamic-range-fraction (DRF) calibration framework. Our quantile-matching calibration is post-hoc, model-agnostic, and

applied directly to VCM outputs without retraining. We note that causal F1 against a fixed real-data GIES ground truth is rank-based on edge sets and is therefore not subject to the same metric-calibration concerns [Miller et al. \(2025\)](#) raise for delta-correlation and per-gene MSE. Our finding that ElasticNet matches or exceeds deep models on causal F1 is therefore complementary to, rather than contradicted by, the DRF results: prediction calibration and causal-recovery calibration are different evaluation axes.

3. Methods

3.1. Datasets

We use the K562 and RPE1 Perturb-seq datasets from [Replogle et al. \(2022\)](#) (GEO accession GSE221321). The K562 dataset comprises 310,385 single-cell transcriptomes spanning 1,861 perturbation-eligible CRISPRi targets (≥ 10 cells per perturbation) and 10,691 non-targeting control cells across 8,563 genes. RPE1 comprises 247,914 cells and 2,069 perturbation-eligible targets. The Norman 2019 CRISPRa dataset ([Norman et al., 2019](#)) (GEO accession GSE133344) provides a cross-paradigm test: 111,445 K562 cells under 237 perturbations, filtered to 102 single-gene perturbation targets after removing combinatorial conditions. For tractability, primary evaluations use a 200-gene subset selected with numpy random seed 42 from the K562 perturbation-eligible set; the Norman evaluation uses all 102 eligible single-gene targets. Scale-sweep evaluations use $N \in \{50, 75, 100, 125, 150, 175, 200\}$.

3.2. Scope of empirical claims

To pre-empt scope ambiguity, Table 1 enumerates exactly which combinations of dataset, predictive method, causal-discovery backend, and seed budget produced the numerical results reported in this paper. Empirical claims outside this table are explicit limitations (§5) rather than findings.

Dataset	Predictive methods
Replogle K562	ENet, CPA, GEARS, Geneformer, STATE, GRNBoost2, overlay-real
Replogle RPE1	ENet, GEARS, CPA, GRNBoost2, overlay-real
Norman 2019	ENet, GEARS, CPA, GRNBoost2, overlay-real

Table 1. Scope-of-empirical-claims table. Deep-model substitutability is tested with two architecturally distinct trained VCMs (GEARS, CPA) on the cross-context (RPE1) and cross-paradigm (Norman) datasets in addition to K562; both recover no more than chance on both. Geneformer multi-dataset is omitted (its prediction is a similarity heuristic, not learned ISP expression, §5).

3.3. Virtual cell models

Three deep VCMs are evaluated using their canonical architectures and default hyperparameters:

- GEARS ([Roohani et al., 2024](#)): graph neural network with Gene Ontology gene-coexpression knowledge graph; trained 15 epochs on Apple M4 CPU; validation MSE 3.69; ~ 70 min wall time.
- CPA ([Lotfollahi et al., 2023](#)): compositional perturbation autoencoder (conditional VAE) with adversarial covariate disentanglement; trained 50 epochs on Apple M4 CPU; validation $r^2(\text{mean}) = 0.71$; ~ 6 min wall time.
- Geneformer V2 ([Theodoris et al., 2023](#)): 104M-parameter transformer pretrained on ~ 30 M single-cell profiles; in-silico perturbation via embedding cosine similarity; ~ 9 min on Apple M4 CPU.

All other hyperparameters use package defaults at versions specified in Section 3.10.

3.4. ElasticNet baseline

For each of the 200 evaluation genes, an `sklearn.linear_model.ElasticNet` is fit on held-out non-targeting control cells ($\alpha = 0.1$, `l1_ratio` = 0.5, `max_iter` = 1000, `random_state` = 0). At prediction time, the perturbed gene is set to 10% of its control mean (modeling $\sim 90\%$ CRISPRi knockdown) and the 199 remaining ElasticNets predict downstream gene expression. Real perturbation values are not used at training or prediction time, eliminating leakage ([Zou & Hastie, 2005](#)).

3.5. Overlay sampling for covariance preservation

Independent negative binomial sampling from predicted means destroys gene-gene covariance and reduces GIES edge recovery to under 2% of the real-data ceiling. We use overlay sampling: for each perturbation, 100 cells are bootstrap-sampled from

a subsampled control matrix ($n = 200$, matching the CausalBench protocol; [Chevalley et al., 2025](#)), each cell’s gene expression is multiplied element-wise by $(\mu_{\text{pred}} + 0.01)/(\mu_{\text{ctrl}} + 0.01)$. GIES is clipped to $[0.01, 10]$, and

Poisson noise is applied. Sensitivity analysis sweeping clip range, pseudocount, and noise model confirms model rankings are stable; only absolute values shift.

3.6. Causal discovery backends

GIES (primary). Greedy Interventional Equivalence Search ([Hauser & Bühlmann, 2012](#)) via the gies

Python package, partitioned into ten 20-gene partitions per the CausalBench protocol (Chevalley et al., 2025), executed in parallel via Python multiprocessing with `OMP_NUM_THREADS=1`.

inspre (secondary, signal probe). Sparse inverse of the average causal effect matrix

$$\text{ACE}_{ij} = \mathbb{E}[X_j \mid \text{do}(X_i)] - \mathbb{E}[X_j \mid \text{ctrl}]$$

via the official R inspre package called as a Python subprocess (Brown et al., 2025). Regularization parameter selected by held-out test error.

3.7. Cross-paradigm replication on Norman 2019

To test whether the substitutability conclusion generalizes across perturbation modality (CRISPRa rather than CRISPRi) and laboratory (Norman et al. 2019 vs. Replogle et al. 2022), we evaluated the full method panel on Norman 2019 K562 data, restricted to $N=102$ single-gene perturbation-eligible targets (combinatorial conditions dropped so the on-target-edge metric matches K562/RPE1) on an 8,028-gene HVG background chosen for CPA-training tractability. We retrained both GEARS and CPA from scratch on this Norman subset and pushed their predictions through the identical GIES pipeline. Against the real-data Norman ground truth (209 edges), GEARS recovers $F1 = 0.211$ [0.176, 0.248] and CPA $F1 = 0.182$ [0.155, 0.212], both indistinguishable from the random-edge floor ($F1 = 0.179$ [0.147, 0.209]) and the ElasticNet baseline ($F1 = 0.190$ [0.158, 0.223]). The substitutability gap therefore persists under a change of perturbation paradigm: deep VCMs retrained on CRISPRa data recover no more of the real causal graph than chance, while the sparse ElasticNet baseline remains competitive — confirming the conclusion is neither Replogle-K562-specific nor an artifact of the CRISPRi modality.

3.8. Effect-size calibration

Quantile-matching post-hoc calibration is implemented in `causalcellbench/calibration/` using the rank-preserving distribution-mapping algorithm of Bolstad et al. (2003) and Reinhard et al. (2001). For each model:

1. Compute $\log_2 \text{FC}_{ij} = \log_2((\mu_{\text{pred},j}^{(i)} + 0.01)/(\mu_{\text{ctrl},j} + 0.01))$ per (perturbation, gene) entry.
2. Take absolute values, store signs.
3. Compute empirical CDF of $|\log_2 \text{FC}|$ via rank-then-divide-by- n .

4. Look up the corresponding quantile of the real-data distribution via `numpy.interp` over 1,000 quantile knots.

5. Recover calibrated expression as $(\mu_{\text{ctrl}} + 0.01) \cdot 2^{\text{sign} \cdot |\log_2 \text{FC}_{\text{cal}}|} - 0.01$, clipped non-negative.

Calibration parameters are estimated once on the full real-data $|\log_2 \text{FC}|$ distribution from the K562 200-gene perturbation-eligible set (40,000 (perturbation, gene) values).

3.9. Evaluation

Primary metrics are F1, precision, and recall against the real-data GIES edge set (substitutability ground truth):

$$F1 = \frac{2 \cdot \text{TP}}{2 \cdot \text{TP} + \text{FP} + \text{FN}}.$$

Secondary: top- k AUPRC sweep, P@100, Structural Hamming Distance (SHD), mean Wasserstein distance per perturbation, False Omission Rate via Mann-Whitney U with global Benjamini-Hochberg correction. Multi-hop credit at $k=2,3$ implemented via `networkx.has_path`.

Statistical analysis. Headline Table 2 reports across-seed mean and 95% t -distribution CIs from $n = 3$ random seeds (42, 123, 456). Each seed re-randomizes gene subset selection, control cell subsampling, and GIES partition order. Pairwise paired permutation tests (500 permutations) on F1 differences are computed between every pair of predictive methods; significance is reported at $p < 0.05$ uncorrected and Bonferroni-corrected over 21 pairs. The technical-duplicate noise floor (Section 4.4) is computed as the mean F1 between two independent GIES recoveries on disjoint random cell halves of the real K562 dataset, averaged across 3 split seeds. Bootstrap 200-iteration edge-set CIs are reported per seed in the released JSONs but not in the headline table to avoid double-reporting variance.

3.10. Implementation

Python 3.10 with `anndata` 0.10, `scanpy` 1.10, `gies` 0.1, `arboreto` 0.1.6, `scikit-learn` 1.5, `networkx` 3.3. R 4.5.2 with `inspre` 1.0.1 called via subprocess. Hardware: Apple M4 (local) for all reported numbers. All code, configurations, and edge caches are released under MIT license at an anonymized mirror linked from the supplementary submission portal; the de-anonymized GitHub repository will be released at camera-ready time.

4. Results

4.1. Causal recovery on K562 at $N=200$

We evaluate all conditions across five random seeds (42, 123, 456, 789, 1024), each re-randomizing gene subset selection, control cell subsampling, and GIES partition order. Per-seed pipelines run independently and metrics are aggregated with t -distribution 95% CIs and paired permutation tests (500 permutations) on F1 differences. Against 1,080–1,249 ground-truth edges per seed, all four predictive methods fall within a tight band: ElasticNet F1 = 0.426 [0.406, 0.446], CPA = 0.412 [0.392, 0.432], GEARS = 0.406 [0.388, 0.424], Geneformer = 0.425 [0.412, 0.439], overlay-resampled real = 0.422 [0.412, 0.433] (Table 2; Figure 1). No pairwise difference among predictive methods is significant at $p < 0.05$ across seeds; ElasticNet (0.426) and Geneformer (0.425) are statistically tied. The observational baseline GRNBoost2 falls to F1 0.098 [0.094, 0.103], with 95% CIs disjoint from every predictive method by ≥ 0.27 F1; with $n = 5$ paired sign-flip permutation, the floor two-sided p -value is $2/2^5 = 0.0625$, so $p < 0.05$ is not certifiable at this seed budget despite the gap being far larger than the per-method standard error. The random-bootstrap floor is F1 ≈ 0.04 . The technical-duplicate noise floor from cell-split sampling (Section 4.4) is 0.045 ± 0.003 . The headline finding is therefore not that ElasticNet outperforms deep models—the previous single-seed claim does not survive multi-seed evaluation—but that a 104M-parameter pretrained transformer, a graph neural network with Gene Ontology priors, and a conditional VAE with adversarial covariate disentanglement all achieve causal-recovery F1 statistically indistinguishable from a sparse linear regression with $\approx 200^2$ free coefficients.

4.2. Effect-size inflation across architectures

Real K562 perturbation responses at the full transcriptome scope are sparse: 2.7% of (perturbation, gene) effects exceed $|\log_2 \text{FC}| = 0.5$, and the leading singular vector of the perturbation-response matrix captures 56.8% of total variance. All three deep models produce predictions in which the corresponding fraction is 63.2% (GEARS), 63.4% (CPA), and 63.5% (Geneformer)—a 23-fold inflation that is identical across architectures to within 0.3 percentage points (Figure 2). Mean per-perturbation L_1 norm of predicted shifts is approximately tenfold larger than empirical: real K562 mean $L_1 = 22.8$; GEARS = 226.9; CPA = 226.6; Geneformer = 227.4.

The control-regime covariance structure is preserved well by all three models (correlation of correlation ma-

trices: GEARS = 0.967, CPA = 0.997, Geneformer = 0.996), indicating that deep models have learned baseline gene-gene relationships in unperturbed cells but fail to predict sparse, structured deviations from baseline upon perturbation.

4.3. Reconciliation: mode collapse and inflation are two views of one failure

The inflation observation above appears to contradict the recent mode-collapse literature (Mejia et al., 2025; Ahlmann-Eltze & Huber, 2025), which reports that scRNA-seq perturbation models compress per-perturbation variance toward the dataset mean. We test the two against each other on identical predictions and find they coexist as two views of the same underlying failure. Per-perturbation Δ -variance ratios (predicted shift variance over real shift variance) are below 1 for all three deep models on the 200-gene scope: median 0.461 (GEARS), 0.432 (CPA), 0.008 (Geneformer); GEARS and CPA compress real per-perturbation variance by approximately half, Geneformer by two orders of magnitude. Inter-perturbation cosine similarity of shift vectors is 0.084 in real data and 0.845 in CPA predictions—CPA produces nearly the same response vector for every perturbation. When per-perturbation effects are pooled across all (perturbation, gene) entries at the full transcriptome scope, the same homogenization manifests as the 23-fold inflation: dense responses at off-target genes accumulate. Mode collapse and inflation are therefore not contradictory; they are the per-perturbation and pooled views, respectively, of a model class that predicts a single dense response pattern across all perturbations.

4.4. Technical-duplicate noise floor

To anchor model F1 values against the intrinsic technical-duplicate ceiling of the dataset (Mejia et al., 2025; Miller et al., 2025), we compute F1 between two GIES recoveries on disjoint random halves of real K562 cells across three random splits. The mean F1 between halves is 0.045 ± 0.003 , with both halves recovering 1,080–1,110 edges and only ~ 50 in common. Real-cell sampling variance therefore dominates GIES at $N=200$ with the cells-per-perturbation budgets used here, and all evaluated methods (ElasticNet 0.425, deep models 0.389–0.432) operate roughly an order of magnitude above this raw cell-split floor. We report this technical-duplicate F1 as a row in Table 2 alongside the random-bootstrap floor (F1 ≈ 0.04) for clarity on what the ranking is and is not measuring: methods are ranked against the GIES-on-real ground truth derived from overlay-resampled cells, not against a noise-

Table 2. Causal recovery on K562 at $N=200$ genes against real-data GIES ground truth. Multi-seed mean F1, precision, recall, and edge count across $n = 5$ random seeds (42, 123, 456, 789, 1024) with 95% t -distribution CIs. “Tech. duplicate” is F1 between two GIES recoveries on disjoint random halves of real cells (mean of 3 splits). Pairwise paired-sign-flip permutation tests on F1 differences find no significant difference at $p < 0.05$ among the four predictive methods (ElasticNet, CPA, GEARS, Geneformer). GRNBoost2 sits at F1 0.098 with 95% CI disjoint from every predictive method by ≥ 0.27 F1 (well over 10 standard errors); however, the $n = 5$ paired-sign-flip permutation distribution has a minimum two-sided p -value of $2/2^5 = 0.0625$, so $p < 0.05$ is not certifiable at this seed budget despite the obvious effect-size gap. ElasticNet (0.426) and Geneformer (0.425) are statistically tied.

Condition	F1 [95% CI]	Precision [95% CI]	Recall [95% CI]	Edges
Real Data (ceiling)	1.000	1.000	1.000	—
Overlay Real (UB, $n=20$)	0.429 [0.419, 0.438]	0.418 [0.409, 0.427]	0.440 [0.429, 0.451]	1,076
ElasticNet ($n=20$)	0.418 [0.408, 0.427]	0.406 [0.396, 0.416]	0.430 [0.420, 0.441]	1,076
CPA ($n=5$)	0.412 [0.392, 0.432]	0.407 [0.390, 0.423]	0.417 [0.393, 0.441]	1,094
GEARS ($n=5$)	0.406 [0.388, 0.424]	0.404 [0.388, 0.419]	0.408 [0.385, 0.431]	1,099
Geneformer ($n=5$)	0.425 [0.412, 0.439]	0.422 [0.409, 0.434]	0.429 [0.413, 0.445]	1,071
GRNBoost2 (obs, $n=20$)	0.098 [0.096, 0.101]	0.074 [0.072, 0.076]	0.147 [0.144, 0.151]	3,119
Tech. duplicate (noise floor)	0.045 [0.042, 0.048]	0.045	0.045	$\sim 1,090$
Random bootstrap (floor)	≈ 0.04	—	—	—

less oracle.

4.5. Inspre signal probe and full lambda-path validation

The inspire algorithm (Brown et al., 2025), applied to the ACE matrix from each condition without partitioning, recovers 12,280 edges from real Perturb-seq, exactly 12,280 from ElasticNet predictions, 6,234 from GEARS predictions, and zero edges from both CPA and Geneformer at the test-error-optimal regularization parameter (Figure 3). To address whether the zero result is a binary threshold artifact or absolute, we trace recovered edge count across the full λ path of inspire’s coordinate-descent solver for each condition. GEARS recovers up to 6,634 edges across the path (smallest λ with any edge: 5.78). CPA’s vanilla path yields no finite λ values (the inspire solver fails on its ACE matrix), supporting the strongest reading: CPA’s predicted ACE matrix has signal-to-noise too low for sparse inverse regression to even traverse a valid regularization path. Geneformer recovers up to 602 edges (smallest λ with any edge: 0.67), so the zero-edges result for vanilla Geneformer at the test-error-optimal λ does not generalize to the entire path. Quantile calibration substantially expands the lambda range over which edges are detectable: calibrated CPA recovers up to 419 edges (smallest λ : 0.20) where vanilla CPA recovered none, and calibrated Geneformer recovers up to 1,315 edges (smallest λ : 0.014) where vanilla recovered 602. The GIES-side calibration result (Geneformer F1 -0.009) and the inspire-side calibration result (Geneformer $+713$ max edges along path) are not contradictory: GIES partitioning forces equal-cardinality outputs, so a Geneformer prediction set that is more diffuse but contains

more inspire-detectable signal can lose F1 against a fixed-edge-count GIES baseline while gaining recoverable signal in inspire.

4.6. Scale sweep across $N=50-200$

The deep model achieving the highest precision rotates across gene-set sizes: Geneformer leads at $N=150$ (precision 0.394), GEARS at $N=175$ (0.401), and ElasticNet at $N=200$ (0.378). ElasticNet is in the top two methods at every scale tested. No deep model dominates consistently. Single-scale evaluation of these methods would yield three different headline conclusions depending on chosen N .

4.7. Reversed-direction failure mode

Three pieces of evidence point to a single conclusion: deep VCM predictions miss the directionality of causal regulation more than they miss the regulatory pairs themselves. First, multi-hop credit analysis (a predicted edge is correct at hop- k if a path of length $\leq k$ exists in the ground-truth graph) increases precision by 9.2 pts (Geneformer) to 11.8 pts (GEARS) at $k=2$ across all conditions. Second, the SHD breakdown across all four deep models shows 27–30% of false-positive predictions are gene pairs connected in the ground-truth graph but with direction reversed (Table 2 “Reversed” column; multi-seed JSON in supplementary). Third, undirected F1 (treating both predicted and ground-truth edges as unordered pairs) lifts every method’s score by approximately 0.06–0.09 F1 over directed F1, recovering most but not all of the reversed-edge gap. We note that GIES emits CPDAGs (partially directed graphs) and that the directed F1 reported throughout is computed after canonicalizing

Table 3. Vanilla vs. calibrated head-to-head on K562 at $N=200$, identical seed and partition configuration. $\Delta F1$ is calibrated minus vanilla. MSE Δ is reconstruction MSE percentage change.

Model	Vanilla F1	Cal. F1	$\Delta F1$	MSE Δ
CPA	0.3887	0.4051	+0.016	+6%
GEARS	0.4055	0.4133	+0.008	+33%
Geneformer	0.4324	0.4232	-0.009	+8%

undirected CPDAG arrows by gene-index ordering; the undirected F1 sanity check confirms that the headline tied-deep-model finding is not an artifact of this canonicalization (the tie is observable on the undirected representation as well).

4.8. Cross-context replication on RPE1

The structural properties of RPE1 perturbation responses differ from K562: leading singular variance 34.5% (vs. 56.8%), fraction near-zero 15.4% (vs. 38.6%), fraction large 17.3% (vs. 2.7%). We retrained both GEARS and CPA from scratch on RPE1 (8,749 genes, 196 perturbation-eligible targets) and pushed their predictions through the identical GIES pipeline. Against the real-data RPE1 ground truth (1,137 edges), GEARS recovers $F1 = 0.321$ [0.304, 0.337] and CPA $F1 = 0.361$ [0.343, 0.377] — both statistically indistinguishable from the random-edge floor ($F1 = 0.347$ [0.329, 0.363]) and the ElasticNet baseline ($F1 = 0.340$ [0.324, 0.358]); all bootstrap confidence intervals overlap. The substitutability gap therefore persists under a change of cellular context: two architecturally distinct deep VCMs retrained on RPE1 recover no more of the real causal graph than a sparse linear baseline or chance.

4.9. Effect-size calibration

The quantile-matching calibration corrects the magnitude distribution exactly: post-calibration $|\log_2 FC| > 0.5$ fraction is 11.2% (GEARS), 11.4% (CPA), 11.4% (Geneformer), within 0.2 pts of the real-data 200-gene reference 11.4% (Figure 4, top). Downstream causal F1 improves modestly for two of three models: CPA $0.389 \rightarrow 0.405$ ($\Delta = +0.016$), GEARS $0.406 \rightarrow 0.413$ ($\Delta = +0.008$), Geneformer $0.432 \rightarrow 0.423$ ($\Delta = -0.009$) (Figure 4, bottom; Table 3). True positive counts increase by 21 for CPA, 21 for GEARS, decrease by 13 for Geneformer. Reconstruction MSE increases modestly: CPA +6%, GEARS +33%, Geneformer +8%. No calibrated deep model surpasses ElasticNet ($F1 = 0.425$).




Figure 1. Causal recovery on K562 at $N=200$ genes against the real-data GIES ground truth, aggregated across five random seeds (42, 123, 456, 789, 1024). F1 score for each evaluated condition: real Perturb-seq (substitutability ceiling); ElasticNet sparse linear regression baseline (Zou & Hastie, 2005); overlay-resampled real perturbation means (upper bound for any mean-prediction model under the same overlay sampling protocol); three deep VCMs (GEARS, CPA, Geneformer); and an observational baseline (GRNBoost2). Bars indicate across-seed mean F1; error bars are 95% t -distribution CIs. The red dashed line marks the technical-duplicate noise floor (mean F1 = 0.045 between two GIES recoveries on disjoint random halves of real cells, mean of three splits). The dotted line marks the random-bootstrap floor.

5. Discussion

Why does the failure replicate across architectures? All three deep VCMs produce similar mode-collapse signatures: per-perturbation Δ -variance ratios below 1, inter-perturbation cosine similarities elevated above the real-data 0.087, and pooled $|\log_2 FC|$ inflation at the full-transcriptome scope. GEARS, CPA, and Geneformer differ in architecture, training objective formulation, and parameter count by orders of magnitude—so the shared signature is suggestive of an objective-driven rather than architecture-driven cause. To test this, we retrained CPA with Gaussian (textbook per-gene MSE) reconstruction loss in place of its default Negative Binomial loss, holding all other hyperparameters constant; if the failure were purely an artifact of the NB reconstruction term, the alternative loss would reduce mode collapse and improve causal recovery. The result is striking: the Gaussian-loss variant produces more per-perturbation variance compression than vanilla CPA (median Δ -variance ratio $0.43 \rightarrow 0.18$) and only marginally less inter-perturbation similarity (cosine $0.84 \rightarrow 0.80$), yet higher causal F1 against the GIES ground truth ($0.389 \rightarrow 0.414$ at seed 42). Per-perturbation magnitude and causal recovery




Figure 2. Mode-collapse vs. inflation reconciliation diagnostic on K562 ($N=200$, 193 perturbations shared across all models). (a) Distribution of per-perturbation Δ -variance ratios (predicted shift variance / real shift variance) per model; the vertical dashed line marks no compression (ratio = 1). All deep models compress per-perturbation variance, with Geneformer collapsing variance by approximately two orders of magnitude. (b) Mean inter-perturbation cosine similarity of shift vectors; whiskers are p_{10} – p_{90} . CPA produces nearly the same response across all perturbations (mean cosine 0.84 vs. 0.08 for real). (c) Mean per-perturbation fraction $|\log_2 \text{FC}| > \tau$ as a function of threshold τ ; deep models track real or fall below at all thresholds when computed per-perturbation, although they exceed real at the full-transcriptome pool of (perturbation, gene) entries. $n = 200$ genes $\times$ 193 perturbations per condition.

are dissociated: a model can produce more aggressively compressed per-perturbation predictions and recover more of the real-data causal graph. The shared failure is therefore neither in the reconstruction loss form nor in the marginal magnitude distribution; it is in the joint covariance and conditional-independence structure that GIES extracts from the multi-cell intervention regime. A direct test: train CPA with an explicit Frobenius-norm penalty on the discrepancy between predicted and real perturbation-response covariance, holding the loss family constant, and measure whether the resulting model exceeds the alt-loss F1 of 0.414.

Architectural ablation across six CPA variants. To extend the alt-loss probe into a full architectural ablation, we trained six CPA variants—each modifying exactly one architectural component relative to the NB-loss baseline—and ran the identical overlay-sampling and GIES pipeline (seed 42, partition size 20) on the predictions of each. The variants cover the four architectural axes most plausibly responsible for joint-distribution distortion: loss family (loss_gauss: Gaussian instead of Negative Binomial), latent capacity (latent_16 and latent_256 vs. baseline 64), decoder depth (decoder_1layer vs. baseline 3), normalization (layer_norm replacing batch-norm in the decoder), and adversarial covariate disentanglement (no_adversarial with reg_adv = pen_adv = 0). Per-variant causal F1 against the real-data GIES ground truth (1,220 edges): latent_256 0.418, decoder_1layer 0.418, latent_16 0.413, loss_gauss 0.401, layer_norm 0.399, no_adversarial 0.387. The full ablation range is F1 0.387–0.418 (spread 0.031), which sits within the

95% CI of vanilla CPA (0.412 [0.392, 0.432]). No single architectural component, when modified in isolation, moves causal recovery F1 outside the noise envelope of the baseline. The component contributing most to F1 is the adversarial covariate disentangler (removing it costs 0.025 F1, the largest single drop), but removing it does not reproduce the gap to real-data GIES recovery (F1 = 1.000 by construction). The collective implication is that the $\sim 60\%$ of substitutability gap that magnitude calibration leaves uncorrected is not localizable to any one architectural component of CPA either; the gap lives in joint-distribution structure that survives loss/capacity/depth/normalization/adversarial modifications.

Why does ElasticNet match (not beat) deep models? Across three seeds, ElasticNet F1 (0.416) and the three deep models (0.405–0.425) fall within overlapping 95% CIs. The honest reading is that a sparse linear regression with $\approx 200^2$ free coefficients trained only on control cells matches a 104M-parameter pretrained transformer on causal recovery; it does not significantly outperform it. We attribute this to inductive-bias matching: L_1 regularization in ElasticNet structurally biases predictions toward sparsity (Zou & Hastie, 2005), which matches the empirical sparsity of real perturbation responses (38.6% of K562 effects near zero, 56.8% of variance in the leading singular vector). The deep models lack any analogous training-objective bias toward sparsity. A direct test: introduce an L_1 -inducing structured-sparsity penalty into CPA’s training, re-train, and compare the resulting per-perturbation Δ -variance ratio (Section 4.3) and causal F1 against




Figure 3. Inspre signal probe across data sources at the test-error-optimal regularization parameter. Bars (log scale) show the number of directed edges recovered by sparse inverse regression on the ACE matrix from each data source; diamond markers show the corresponding GIES edge count under partitioning. Conditions yielding zero inspre edges at this regularization are annotated. The full- λ -path validation reported in Section 4 shows that vanilla Geneformer and calibrated CPA recover edges at smaller-than-cross-validated regularizations (Geneformer up to 602; calibrated CPA up to 419), so the zero-edges result is binary at the cross-validated λ but continuous across the path. CPA vanilla is the only condition where the inspre solver fails to traverse a valid λ path, supporting the strongest signal-failure interpretation for that model. $n = 200$ perturbations $\times$ 200 genes per ACE matrix.

vanilla CPA.

Algorithmic robustness via PC and NOTEARS. Adding two algorithmically distinct backends—PC (constraint-based; Fisher-Z, $\alpha=0.05$; [Spirtes & Glymour 2000](#)) and NOTEARS (continuous-optimization with acyclicity constraint; [Zheng et al. 2018](#))—confirms the qualitative ranking. Both yield sparser real-data graphs than GIES (PC: 682, NOTEARS: 328 vs. GIES: 1,220 edges) and per-model F1s near zero for all calibrated variants (NOTEARS) and a CPA/GEARS rank reversal between backends (full table in Appendix ??). The four-backend evidence supports the qualitative claim that no deep predictor significantly outperforms ElasticNet across score-, sparse-inverse-, gradient-, and constraint-based causal discovery, without committing to a single ground-truth ranking from any one backend.

What does the inspre lambda-path result mean? The full-path validation tempers the binary “zero edges from CPA and Geneformer” result into a continuous gradient: CPA vanilla still fails to traverse any valid λ , but Geneformer recovers up to 602 (vanilla) / 1,315




Figure 4. Effect-size calibration corrects magnitude distribution exactly but only partially recovers causal F1. (a) Fraction of $|\log_2 FC| > 0.5$ before (hatched) and after (solid) quantile-matching calibration for each deep model on the K562 200-gene scope; the dashed horizontal line marks the real-data 200-gene reference. (b) Causal F1 against the real-data GIES ground truth before (circles) and after (squares) calibration on the identical pipeline (seed 42, partition size 20); arrows connect paired vanilla and calibrated values. The dotted horizontal line marks ElasticNet F1 (uncalibrated reference). $n = 1$ seed $\times$ 200 perturbations $\times$ 200 genes.

(calibrated) edges at non-cross-validated λ settings. Calibration expands the detectable-signal range for both even when it hurts GIES F1 for Geneformer, indicating GIES (top-edge accuracy) and inspre (total recoverable signal) probe different prediction properties; calibration moves the latter further (Appendix B.2).

Why does calibration help two of three models but hurt the third? CPA and GEARS gain 1–2 F1 points after quantile matching erases their magnitude inflation. Geneformer (best vanilla F1 at 0.432) loses 0.9 F1 under the same calibration, plausibly because its foundation-model pretraining yields a magnitude distribution that is already approximately correct; uniform global matching erases this implicit calibration without replacing it. This argues for a learned per-model shrinkage rather than uniform application.

Three-mode calibration sweep with proper train/test discipline. We extend the original Bolstad–Reinhard quantile-matching calibration with two additional modes that target joint-distribution structure: covariance preservation (Mode 2: project predicted shift onto the top-50 principal-subspace of real-perturbation covariance, shrinkage 0.5) and sparsity injection (Mode 3: soft-threshold $|\log_2 FC|$ to match the empirical near-zero fraction). All six permutations of the three modes are evaluated, fit on a held-out 160-perturbation training set and applied to a 40-perturbation test set, with downstream GIES on each calibrated prediction. Table 4 reports the F1 of the best-ordering per model and the corresponding vanilla (uncalibrated) F1 on the same 40-perturbation test split. Across nine ordering choices (three single-mode plus six pipeline permutations), no calibration

ordering exceeds the uncalibrated baseline on any model: GEARS vanilla 0.303 vs. best calibrated 0.289 (-0.014); CPA 0.309 vs. 0.302 (-0.007); Geneformer 0.319 vs. 0.290 (-0.029). With the train/test split applied honestly, the magnitude calibration story does not survive on a held-out perturbation set: post-hoc calibration as currently implemented redistributes prediction mass on the magnitude marginal in a way that hurts, rather than helps, downstream causal recovery. The substitutability gap is therefore not in the magnitude marginal; it is in the joint structure, and joint-structure calibration via the covariance and sparsity modes evaluated here also fails to recover the gap. This is a stronger negative result than the V2.15 single-seed in-domain finding; it implies that closing the substitutability gap requires training-time intervention (e.g. a covariance-preserving regularizer in the perturbation-prediction loss) rather than post-hoc rescaling.

Model	Vanilla F1	Best calibrated F1
GEARS	0.303	0.289
CPA	0.309	0.302
Geneformer	0.319	0.290

Table 4. Phase 4 calibration sweep on K562 test split (40 of 200 perturbations held out). Per-model best F1 over nine ordering choices vs. the uncalibrated baseline on the same test split. No ordering exceeds vanilla; the substitutability gap does not close under post-hoc magnitude or joint-structure calibration.

STATE foundation model on K562. Routing the K562 200-gene subset through Arc Institute’s SE-600M cell encoder and the ST-SE-Replogle zero-shot/k562 checkpoint (Arc Institute, 2025), STATE recovers $F_1 = 0.462$ (1061 edges; precision = 0.497, recall = 0.432) against real-data GIES ground truth at $N=200$ at seed 42. Taken alone this single seed would be the first model in the roster to materially exceed the ElasticNet baseline ($F_1 \approx 0.32$). It does not survive multi-seed evaluation.

STATE multi-seed: the single-seed result does not replicate. Re-running the identical SE-600M $\rightarrow$ ST-SE-Replogle $\rightarrow$ overlay $\rightarrow$ GIES pipeline at two further seeds (320-cell subsets resampled with seeds 123 and 456) yields $F_1 = 0.040$ and $F_1 = 0.051$ — an order of magnitude below the seed-42 value. The three-seed mean is $F_1 = 0.184$ with a 95% confidence interval of $[-0.088, 0.457]$ that spans zero and sits below the ElasticNet baseline. The seed-42 “narrowing” was a fortunate draw of the 1-cell-per-perturbation sampling rather than a robust property of the model. We therefore cannot conclude that STATE exceeds the linear

baseline: its apparent advantage is within sampling noise, and the central substitutability gap holds for the 2025 foundation model as well. The instability is itself diagnostic — a single perturbed cell per gene is too few to estimate the STATE-predicted perturbation mean stably, so a definitive STATE assessment requires the higher-cell-count embedding that the ~ 28 s/cell SE-600M CPU cost placed outside the 6h Modal budget (Appendix B.2). Absolute F_1 at every seed remains below 0.5.

Limitations. Six limitations are noted. Ground truth. Track A ground truth (real-data GIES) is internally consistent but does not establish biological correctness; a Track B evaluation against the literature-curated TRRUST v2 database (Han et al., 2018) is implemented in the released code but reported as a robustness check (Spearman concordance with Track A model rankings rather than absolute F1 against TRRUST). Pretraining contamination. Geneformer V2’s Genecorpus pretraining includes Replogle 2022, so its Replogle-only is partly memorization — one reason we exclude it from the cross-dataset trained-model comparison above. A clean contamination test would need a temporally held-out post-pretraining dataset: a 2024+ public Perturb-seq atlas on a Geneformer-corpus-excluded human cell type with ≥ 50 perturbation-eligible genes, which is not currently available. Seed budget. Multi-seed CIs reported here use $n \leq 20$ seeds for K562 and $n = 1$ for RPE1 and Norman; the qualitative finding that pairwise differences among predictive methods are not significant is robust to seed count, but minimum detectable F1 differences at $n = 20$ are approximately 0.011. Cross-dataset deep-model scope. Both GEARS and CPA were retrained from scratch and evaluated through the full GIES pipeline on RPE1 (cross-context) and Norman (cross-paradigm); all four model $\times$ dataset cells recover no more of the real causal graph than chance (§3.2), and STATE multi-seed (§5) shows no robust advantage on K562. The cross-context and cross-paradigm claims therefore rest on two architecturally distinct trained deep models, not the ElasticNet baseline alone. Geneformer is omitted on RPE1/Norman because its prediction in our pipeline is a cosine-similarity heuristic over the control mean rather than learned in-silico-perturbation expression, so it would not constitute a trained-model test. Model roster. GEARS, CPA, and Geneformer were chosen for architectural diversity at the time of pipeline construction; they predate STATE (Adduri et al., 2025; Arc Institute, 2025) and other 2025–2026 perturbation models. Recent benchmarks (Wong et al., 2025; Csendes et al., 2025; Wei et al., 2025; Viñas Torné et al., 2025) cover

additional models on the per-gene prediction task but not on the causal-recovery task evaluated here. Geneset scale. Primary results are at $N = 200$ for computability of GIES partitioning; scale-sweep numbers extend to $N = 200$ along seven points and inspire extends scaling to ~ 1000 genes, but a head-to-head deep-model evaluation at $N = 500$ – 1000 is not yet completed.

6. Conclusion

PerturbCausal is the first benchmark testing whether virtual cell model predicted Perturb-seq data can substitute for real Perturb-seq in downstream causal gene regulatory network recovery. Across one primary dataset (Replogle K562) with up to 20 random seeds, five predictive methods (ElasticNet, CPA, GEARS, Geneformer, STATE), and four causal-discovery backends (GIES, inspire, PC, NOTEARS), no deep virtual cell model significantly exceeds a sparse linear regression baseline on causal-recovery F1; STATE’s single-seed apparent narrowing does not replicate across seeds (three-seed $F_1 = 0.184$, CI spanning zero), and CPA retrained on a second cell context (RPE1) and a second perturbation paradigm (Norman CRISPRa) likewise recovers no more than chance. The substitutability failure decomposes into a magnitude marginal (mode collapse + effect-size inflation, the two equivalent views identified by Mejia et al. (2025) and our diagnostic, partially correctable post-hoc via Bolstad–Reinhard quantile matching) and a joint-distribution residual (covariance and edge-orientation structure that distributional calibration cannot reach); the benchmark, three-mode calibration diagnostic, multi-seed pipeline, and reproducibility infrastructure are released as the open-source PerturbCausal Python package.

References

Ahlmann-Eltze, C. and Huber, W. Deep-learning-based gene perturbation effect prediction does not yet outperform simple linear baselines. *Nature Methods*, 2025.

Bolstad, B. M., Irizarry, R. A., Astrand, M., and Speed, T. P. A comparison of normalization methods for high density oligonucleotide array data based on variance and bias. *Bioinformatics*, 19(2):185–193, 2003.

Brown, B. C., Spence, J. P., and Pritchard, J. K. Large-scale causal discovery using interventional data sheds light on gene network structure in K562 cells. *Nature Communications*, 16:9628, 2025.

Bunne, C., Roohani, Y., Rosen, Y., et al. How to build the virtual cell with artificial intelligence: priorities and opportunities. *Cell*, 187:7045–7063, 2024.

Chevalley, M., Roohani, Y., Mehrjou, A., Leskovec, J., and Schwab, P. CausalBench: A large-scale benchmark for network inference from single-cell perturbation data. *Communications Biology*, 2025.

Cui, H., Wang, C., Maan, H., et al. scGPT: toward building a foundation model for single-cell multi-omics using generative AI. *Nature Methods*, 21:1470–1480, 2024.

Han, H., Cho, J.-W., Lee, S., et al. TRRUST v2: an expanded reference database of human and mouse transcriptional regulatory interactions. *Nucleic Acids Research*, 46(D1):D380–D386, 2018.

Hauser, A. and Bühlmann, P. Characterization and greedy learning of interventional Markov equivalence classes of directed acyclic graphs. *Journal of Machine Learning Research*, 13:2409–2464, 2012.

Lotfollahi, M., Klimovskaia Susmelj, A., De Donno, C., et al. Predicting cellular responses to complex perturbations in high-throughput screens. *Molecular Systems Biology*, 19:e11517, 2023.

Mejia, G. M., Kandari, P., Rai, A., et al. Diversity by design: Addressing mode collapse improves scRNA-seq perturbation modeling on well-calibrated metrics. *arXiv preprint arXiv:2506.22641*, 2025.

Miller, H. E., Mejia, G. M., Leblanc, F. J. A., et al. Deep learning-based genetic perturbation models do outperform uninformative baselines on well-calibrated metrics. *bioRxiv* 2025.10.20.683304, 2025.

Reinhard, E., Adhikhmin, M., Gooch, B., and Shirley, P. Color transfer between images. *IEEE Computer Graphics and Applications*, 21(5):34–41, 2001.

Replogle, J. M., Saunders, R. A., Pogson, A. N., et al. Mapping information-rich genotype-phenotype landscapes with genome-scale Perturb-seq. *Cell*, 185(14):2559–2575, 2022.

Roohani, Y., Huang, K., and Leskovec, J. Predicting transcriptional outcomes of novel multigene perturbations with GEARS. *Nature Biotechnology*, 42:927–935, 2024.

Roohani, Y., Lerman, J., and Leskovec, J. Virtual cell challenge: Toward a Turing test for the virtual cell. *Cell*, 188:3370–3382, 2025.

- Theodoris, C. V., Xiao, L., Chopra, A., et al. Transfer learning enables predictions in network biology. *Nature*, 618:616–624, 2023.
- Callahan, T., Bertin, P., Tejada-Lapuerta, A., et al. Virtual cells as causal world models: A perspective on evaluation. In *NeurIPS 2025 AI4D3 Workshop*, 2025. URL: https://ai4d3.github.io/2025/papers/25_Virtual_Cells_as_Causal_Wor.pdf.
- Wu, M., Bao, Y., Xu, K., Patel, A., Bai, K., Schiebinger, G., and Lopez, R. PerturBench: Benchmarking machine learning models for cellular perturbation analysis. *arXiv preprint arXiv:2408.10609*, 2024.
- Wu, M., Tu, R., Lopez, R., Schiebinger, G., and Bai, K. Identifying biological perturbation targets through causal differential networks. In *International Conference on Machine Learning*, 2025.
- Zheng, X., Aragam, B., Ravikumar, P., and Xing, E. P. DAGs with NO TEARS: Continuous optimization for structure learning. In *NeurIPS*, 2018.
- Arc Institute. STATE: Multi-scale AI model for cellular perturbation effects. GitHub: <https://github.com/ArcInstitute/state>; HuggingFace: <https://huggingface.co/arcinstitute/SE-600M>, 2025.
- Adduri, A. et al. Predicting cellular responses to perturbation across diverse contexts with STATE. *bioRxiv*, 2025.
- Wong, D. R. et al. Simple controls exceed best deep learning algorithms and reveal foundation model effectiveness for predicting genetic perturbations. *Bioinformatics*, 2025.
- Csendes, G. et al. Benchmarking foundation cell models for post-perturbation RNA-seq prediction. *BMC Genomics*, 2025.
- Wei, Z. et al. Benchmarking algorithms for generalizable single-cell perturbation response prediction. *Nature Methods*, 2025.
- Viñas Torné, R. et al. Systema: a framework for evaluating genetic perturbation response prediction beyond systematic variation. *Nature Biotechnology*, 2025.
- Bendidi, I. et al. Benchmarking transcriptomics foundation models for perturbation analysis: one PCA still rules them all. *arXiv*, 2024.
- Wenteler, A. et al. PertEval-scFM: Benchmarking single-cell foundation models for perturbation effect prediction. *bioRxiv*, 2025.
- Zhu, D. et al. Gradient-aware modeling advances AI-driven prediction of genetic perturbation effects. *bioRxiv*, 2025.
- Li, L. et al. A systematic comparison of single-cell perturbation response prediction models. *bioRxiv*, 2025.
- Cao, Z.-J. et al. Causality-oriented regulatory inference and rational transcriptomic reprogramming design with CASCADE. *bioRxiv*, 2025.
- Park, K. et al. Confounder-robust causal discovery and inference in Perturb-seq using proxy and instrumental variables. 2026.
- Pratapa, A. et al. Benchmarking algorithms for gene regulatory network inference from single-cell transcriptomic data (BEELINE). *Nature Methods*, 2019.
- Kernfeld, E. et al. Transcriptome data are insufficient to control false discoveries in regulatory network inference. *Cell Systems*, 2024.
- Qiu, X. et al. Inferring causal gene regulatory networks from coupled single-cell expression dynamics using Scribe. *Cell Systems*, 2020.
- Spirtes, P. and Glymour, C. *Causation, Prediction, and Search*, 2nd edition. MIT Press, 2000.
- Zou, H. and Hastie, T. Regularization and variable selection via the elastic net. *Journal of the Royal Statistical Society: Series B*, 67(2):301–320, 2005.
- Norman, T. M., Horlbeck, M. A., Replogle, J. M., Ge, A. Y., Xu, A., Jost, M., Gilbert, L. A., and Weissman, J. S. Exploring genetic interaction manifolds constructed from rich single-cell phenotypes. *Science*, 365(6455):786–793, 2019. doi:10.1126/science.aax4438.

A. Additional results

A.1. Scale sweep precision across $N=50-200$

N (genes)	GT edges	ElasticNet	GEARS	CPA	Geneformer
50	169	0.345	0.280	0.302	0.292
75	309	0.309	0.351	0.344	0.316
100	509	0.344	0.339	0.335	0.348
125	646	0.356	0.360	0.326	0.349
150	821	0.368	0.310	0.347	0.394
175	1,002	0.369	0.401	0.378	0.340
200	1,217	0.378	0.392	0.385	0.373

Table 5. Scale-sweep precision against real-data GIES ground truth at seven gene-set sizes. Bold values indicate top performance per row.

A.2. Multi-hop credit at $k = 1, 2, 3$

Method	P@ $k=1$	P@ $k=2$	P@ $k=3$	Boost ($k=2$)	Boost ($k=3$)
ElasticNet	0.425	0.540	0.543	+11.5	+11.9
GEARS	0.410	0.529	0.533	+11.8	+12.2
CPA	0.387	0.500	0.508	+11.3	+12.0
Geneformer	0.432	0.525	0.529	+9.2	+9.7
Overlay Real	0.423	0.533	0.539	+11.0	+11.6

Table 6. Multi-hop credit precision at hop- k over the $N=200$ K562 evaluation. Boost columns show percentage-point increase from $k=1$.

A.3. RPE1 cross-context structural comparison

Property	K562	RPE1
First singular vector variance	56.8%	34.5%
Top-5 singular variance	69.8%	68.5%
Frac near-zero effects	38.6%	15.4%
Frac large effects ($ \log_2 \text{FC} > 0.5$)	2.7%	17.3%
GIES edges (real data)	1,220	810

Table 7. Structural properties of real perturbation-response matrices for K562 and RPE1 at $N=200$.

A.4. Multi-seed bootstrap confidence

Across five seeds (42, 123, 456, 789, 1024) the four predictive methods sit within overlapping intervals on every metric (Table 8). The GRNBoost2 random-baseline floor confirms that all four predictive methods recover real signal substantially above chance, while none of the deep models significantly separates from ElasticNet at this seed budget.

A.5. Full PC and NOTEARS robustness tables

The main-text PC and NOTEARS Discussion paragraph cites summary statistics. Tables 9 and 10 report the full per-condition F1, precision, recall, and edge count against each backend’s own real-data ground truth. Both backends produce sparser graphs than GIES at $N=200$ (PC: 682 edges, NOTEARS: 328 edges, vs. GIES: 1,220), and per-model F1s are uniformly small in absolute terms. Calibration zeros out NOTEARS recovery for all models because quantile matching pushes ACE entries below the $w_{\text{thresh}}=0.1$ weight threshold; under PC the calibrated and vanilla scores are within ± 0.01 for all three models.

Method	F1	Precision	Recall	AUPRC
ElasticNet	0.426 ± 0.016	0.420 ± 0.014	0.431 ± 0.019	0.425 ± 0.016
CPA	0.412 ± 0.016	0.407 ± 0.013	0.417 ± 0.019	0.400 ± 0.022
GEARS	0.406 ± 0.015	0.404 ± 0.012	0.408 ± 0.018	0.419 ± 0.012
Geneformer	0.425 ± 0.011	0.422 ± 0.010	0.429 ± 0.013	0.422 ± 0.013
Overlay-real (oracle)	0.422 ± 0.009	0.420 ± 0.010	0.425 ± 0.009	0.420 ± 0.013
GRNBoost2 (random)	0.098 ± 0.004	0.074 ± 0.003	0.147 ± 0.005	0.077 ± 0.005

Table 8. Multi-seed mean and standard deviation across $n = 5$ seeds at $N=200$ on K562 against real-data GIES ground truth (1,220 edges per seed). Overlay-real is the same overlay-sampling pipeline applied to real perturbed cell predictions and serves as a noise-floor for the overlay step itself. GRNBoost2 is a structurally-distant baseline (gradient-boosted regression) included as a near-random reference.

Condition	Edges	F1	Precision	Recall
Real (PC ground truth)	682	—	—	—
GEARS-vanilla	556	0.0549	0.0612	0.0499
GEARS-calibrated	652	0.0480	0.0491	0.0469
CPA-vanilla	268	0.0800	0.1418	0.0557
CPA-calibrated	294	0.0779	0.1293	0.0557
Geneformer-vanilla	142	0.0097	0.0282	0.0059
Geneformer-calibrated	240	0.0043	0.0083	0.0029

Table 9. PC algorithm (Fisher-Z conditional independence at $\alpha=0.05$) per-condition results on K562 ACE matrices at $N=200$, seed 42.

A.6. Six-variant CPA architectural ablation

To localize the substitutability gap to a specific architectural component, six CPA variants were trained, each modifying one architectural axis relative to the NB-loss baseline. Per-variant causal F1 is reported in Table 11. Spread across all six variants is 0.387–0.418 ($\Delta = 0.031$ F1), within the multi-seed CI of vanilla CPA. Removing the adversarial covariate disentangler is the single largest deletion (loses 0.025 F1, the largest drop), but it does not reproduce the gap to real-data GIES recovery. The collective implication is that the substitutability gap is not localizable to any one architectural component of CPA.

B. Supplementary methods

B.1. Virtual cell model hyperparameters

All four predictive models were trained on the same K562 200-gene subset using their authors’ default hyperparameter values where possible; deviations are deliberate matches to the published reference configurations of each model.

B.2. STATE inference pipeline

The STATE result was obtained zero-shot from the released ST-SE-Reprogle/zeroshot/k562 checkpoint via the state CLI v0.10.5. The pipeline has two stages: (1) SE-600M cell encoder (state emb transform) maps each input cell to a 2,058-dimensional X_{state} embedding; (2) the ST-SE-Reprogle perturbation model (state tx infer) maps $(X_{\text{state}}, \text{perturbation})$ pairs to predicted gene-expression vectors over the same $N=200$ gene panel.

Reproducing the published checkpoint required diagnosing three distinct framework-versioning bugs:

1. Decoupled head dimension. The released checkpoint encodes a Llama backbone with `hidden_size=328`, `num_attention_heads=12`, and explicit `head_dim=64`, i.e. $328 \neq 12 \times 64 = 768$. Transformers ≥ 5 add a strict `validate_architecture` check that rejects `hidden_size % num_attention_heads != 0` regardless of an explicit `head_dim`; transformers 4.45 accepts decoupled head dimensions but lacks the (2) compatibility we describe next; transformers 4.52.3 supports both. Fix: `pin transformers==4.52.3` after `pip install -e state`.
2. `is_causal` kwarg pass-through. The state package’s `LlamaBidirectionalModel.forward` sets `is_causal=False`

Condition	Edges	F1	Precision	Recall
Real (NOTEARS ground truth)	328	—	—	—
GEARS-vanilla	24	0.0511	0.3750	0.0274
GEARS-calibrated	0	0.0000	0.0000	0.0000
CPA-vanilla	18	0.0405	0.3889	0.0213
CPA-calibrated	0	0.0000	0.0000	0.0000
Geneformer-vanilla	7	0.0119	0.2857	0.0061
Geneformer-calibrated	0	0.0000	0.0000	0.0000

Table 10. NOTEARS (continuous-optimization with acyclicity constraint) per-condition results on K562 at $N=200$, seed 42, $\lambda=0.05$, weight threshold $w_{\text{thresh}}=0.1$.

Variant	Edges	F1	Precision	Recall
loss_gauss (Gaussian, not NB)	1,215	0.401	0.402	0.400
latent_16 (latent dim 16, base 64)	1,188	0.413	0.418	0.407
latent_256 (latent dim 256)	1,200	0.418	0.422	0.415
decoder_1layer (decoder depth 1)	1,207	0.418	0.420	0.416
layer_norm (LN replaces BN)	1,212	0.399	0.402	0.396
no_adversarial (reg=pen=0)	1,207	0.387	0.389	0.385

Table 11. Six-variant CPA architectural ablation on K562 at $N=200$, seed 42, partition size 20, against real-data GIES ground truth (1,220 edges). All other hyperparameters are held at the vanilla CPA baseline ($n_{\text{latent}}=64$, encoder layers = 3, decoder layers = 3, batch-norm, NB reconstruction loss, adversarial covariate disentanglement, $\text{max_epochs} = 300$).

and forwards it to the parent `LlamaModel.forward` via `**flash_attn_kwargs`. `transformers 4.45`’s signature lacks the `**flash_attn_kwargs` catch-all, so the call raises `TypeError: unexpected keyword argument 'is_causal'`. `transformers 4.52.3` introduces the catch-all and accepts the call.

3. `ALL_PARALLEL_STYLES` is `None` on Modal. In `transformers 4.52.3` specifically, `transformers.modeling_utils.ALL_PARALLEL_STYLES` is initialized to `None` in some container configurations (verified locally and on Modal), making `LlamaModel.__init__`’s `post_init` hit if `v` not in `ALL_PARALLEL_STYLES` and raise `TypeError: argument of type 'NoneType' is not iterable`. Fix: during image build, replace the offending line in `transformers/modeling_utils.py` with `if ALL_PARALLEL_STYLES is not None and v not in ALL_PARALLEL_STYLES::` this propagates to the state CLI subprocess, where a `sitecustomize.py` monkey-patch did not reliably fire.

The full set of attempted patches (14 attempts) and the final working Modal image specification is in `scripts/modal_run_state_v3.py`.

Subset rationale. STATE inference was run on a 320-cell subset of K562: 120 control cells (sampled from the $\sim 10,691$ non-targeting cells) plus 200 perturbed cells (one per gene in the $N=200$ panel). The reduction is required because SE-600M wall-clock inference on Modal CPU instances measures ~ 28 s/cell at `batch_size=8`; the full 38,979-cell K562 panel would take ~ 300 h, which exceeds the 6-hour Modal function timeout. The 320-cell subset embeds in ~ 13.5 minutes and produces 200 distinct STATE-predicted perturbation profiles—one per perturbation, the unit consumed by overlay sampling and downstream GIES.

B.3. Norman 2019 cross-paradigm pipeline

The Norman 2019 (Norman et al., 2019) dataset (GSE133344) was downloaded from CellxGene Census and contains 111,445 K562 cells under 237 unique perturbations (single and combinatorial). After dropping combinatorial ('A+B') perturbations and keeping only control ($n=11,855$ cells) plus single-gene perturbations with at least 10 cells, 102 eligible target genes remain, all measurable in the standard expression panel.

The substitutability evaluation re-uses the K562 protocol unchanged: 200 control cells + 100 cells per perturbation overlay-sampled from the ElasticNet predictions, vstacked into one expression matrix and fed to GIES with partition size 20 and seed 42.

Two practical engineering issues are noted. (i) `multiprocessing.Pool(fork)` for parallel partition execution dead-

Hyperparameter	Value
CPA (Lotfollahi et al., 2023)	
n_latent	64
n_layers_encoder	3
n_layers_decoder	3
batch_size	512
max_epochs	300
n_epochs_kl_warmup	20
n_epochs_adv_warmup	50
n_epochs_pretrain_ae	10
reconstruction loss	Negative Binomial
GEARS (Roohani et al., 2024)	
hidden_size	64
epochs	20
batch_size	32
Geneformer V2 (Theodoris et al., 2023)	
backbone	ctheodoris/Geneformer (HF)
fine-tuning task	in-silico perturbation
forward_batch_size	16
ElasticNet baseline (Zou & Hastie, 2005)	
alpha	0.1
l1_ratio	0.5
max_iter	1,000
STATE (Arc Institute, 2025)	
embedding model	SE-600M
prediction model	ST-SE-Replogle / zeroshot/k562
hidden_size	328
num_attention_heads	12
head_dim	64 (decoupled, $\neq H/h$)
num_hidden_layers	8
intermediate_size	3,072
max_position_embeddings	64
embedding dim (X_{state})	2,058

Table 12. Training hyperparameters per virtual cell model. Geneformer is fine-tuned from the public ctheodoris/Geneformer V2 checkpoint; STATE is run zero-shot from the released ST-SE-Replogle K562 checkpoint.

locks on the semaphore in Modal’s container runtime; switching to sequential GIES execution removes the deadlock at the cost of single-thread CPU throughput (acceptable: total wall time ≈ 28 s on the M4). (ii) Some Norman partitions contain genes with zero variance under particular interventional regimes, producing $1/\omega = \text{inf}$ in the gauss_int score and an effectively infinite hill-climbing loop in GIES. Adding $\mathcal{N}(0, 10^{-6})$ jitter to the partition expression matrix before scoring prevents the degeneracy without measurably altering recovered edge sets in the well-conditioned partitions.

B.4. Causal-discovery backend implementations

GIES (Hauser & Bühlmann, 2012): Greedy Interventional Equivalence Search via the gies Python package (v0.0.1), score function gauss_int_l0_pen, ~ 20 -gene partitions for tractability of the underlying CPDAG search. Output edges are the union of single-direction CPDAG arrows across partitions; bidirectional CPDAG arrows are kept once with the canonical ordering (lower index first).

inspre (Brown et al., 2025): Sparse-inverse regression on the perturbation-conditional ACE matrix via the inspre R package, full λ -path with leave-one-target-out cross-validation. The Modal R kernel (renv-isolated) is invoked via Rscript subprocess to keep the Python pipeline driver-only.

PC (Spirtes & Glymour, 2000): Constraint-based causal discovery via the causal-learn Python package, Fisher-Z conditional independence at $\alpha = 0.05$. Run on each model’s per-perturbation ACE matrix.

NOTEARS (Zheng et al., 2018): Continuous-optimization with acyclicity constraint, $\lambda = 0.05$, weight threshold

$w_{\text{thresh}} = 0.1$. Implementation: reference notears package on the same per-condition ACE inputs.

Overlay sampling (used by GIES, PC, NOTEARS): for each perturbed condition, 100 synthetic cells are produced by drawing 100 control cells with replacement, multiplying by a per-gene ratio of $(\text{predicted-mean} + 10^{-2})$ over $(\text{control-mean} + 10^{-2})$ clipped to $[0.01, 10]$, and Poisson-resampling to count-like expression. This preserves per-cell control covariance structure while imposing the predicted shift in marginal magnitude.

C. Reproducibility

C.1. Software environment

Component	Version
Python	3.9 (local) / 3.11 (Modal)
PyTorch	2.4.1 (CPU)
transformers	4.52.3 (STATE-pinned)
tokenizers	0.21.x
huggingface_hub	0.30.x
anndata	0.10.9
scanpy	1.10.3
numpy	1.26.4
scipy	1.13.1
scikit-learn	1.x
gies (Python)	0.0.1
causal-learn	0.1.x
notears	reference impl. from (Zheng et al., 2018)
inspre (R)	0.0.1
state (Arc Institute)	0.10.5

Table 13. Software versions used end-to-end. Modal containers pin all dependencies via the image build steps in `scripts/modal_run_*.py`.

C.2. Compute resources

Local pipeline (overlay sampling, GIES, PC, NOTEARS, calibration, multi-seed aggregation, plotting) runs on an Apple M4 laptop with 16 GB unified memory. Wall time: ≈ 45 minutes for the full $N=200$, single-seed K562 evaluation.

Modal cloud is used for the model-training and large-scale inference steps: GEARS, CPA, Geneformer, STATE, and the inspre $N=500$ scale point. Modal H100 instances are reserved for GEARS and Geneformer fine-tuning (≈ 1 h each at $N=200$); A10G for CPA training (≈ 45 minutes); 16-CPU 64 GB instances for inspre and the STATE CPU fallback. The largest single Modal job wall-clock is the SE-600M embedding step in the STATE pipeline at ≈ 3.5 hours for the full embedding, ≈ 13 minutes for the 320-cell subset.

C.3. Random seeds

The single-seed numbers throughout the main text use `seed=42`. Multi-seed CI tables in §A.4 use `{42, 123, 456, 789, 1024}`. Inside the GIES partitioning step, the gene-permutation seed is 42; partition contents are deterministic given the panel and seed.

C.4. End-to-end reproduction commands

All numbers in the main text and tables are reproducible from the released code with `seed=42`.

```
# Calibrate predictions
PYTHONPATH=. python3 scripts/calibrate_predictions.py \
  --diagnostics-out results/calibration_diagnostics.json
```

```
# Vanilla evaluation
```

```
935 PYTHONPATH=. OMP_NUM_THREADS=1 python3 \  
936 scripts/run_eval_local.py \  
937 --data-path data/k562_subset_n200_slim.h5ad \  
938 --pre-subset --n-genes 200 --partition-size 20 \  
939 --n-cells 100 --n-ctrl 200 --seed 42 \  
940 --model-outputs-dir data/model_outputs_real_n200 \  
941 --gies-cache-dir results/gies_cache --skip-random  
942  
943 # Calibrated evaluation  
944 PYTHONPATH=. OMP_NUM_THREADS=1 python3 \  
945 scripts/run_eval_local.py \  
946 --data-path data/k562_subset_n200_slim.h5ad \  
947 --pre-subset --n-genes 200 --partition-size 20 \  
948 --n-cells 100 --n-ctrl 200 --seed 42 \  
949 --model-outputs-dir \  
950 data/model_outputs_real_n200_calibrated \  
951 --gies-cache-dir results/gies_cache_calibrated \  
952 --skip-random  
953  
954 # STATE evaluation (after Modal returns state_predictions.h5ad)  
955 PYTHONPATH=. OMP_NUM_THREADS=1 python3 \  
956 scripts/state_eval_gies.py  
957  
958 # Norman 2019 cross-paradigm  
959 .venv_gears/bin/python scripts/run_norman_local.py  
960  
961  
962  
963  
964  
965  
966  
967  
968  
969  
970  
971  
972  
973  
974  
975  
976  
977  
978  
979  
980  
981  
982  
983  
984  
985  
986  
987  
988  
989
```